

Computer Vision and Image Processing

Gil Ben-Artzi

3


Agenda

- Week 1
 - Image Enhancement: histogram, quantization
- Week 2
 - Filtering: smoothing, median filtering, sharpening
 - Low level detection: Template matching, Edges, Line, Circles
- Week 3
 - Image Blending and Pyramids, Optical Flow

Image Blending

What is Blending?


Merge two image in a seamless way



What is Blending?



Blending: Input Images and Mask



Output: New Image






Image Blending





Image Blending: Challenges






Ghosting






Visible Seams


Ghosting: Window size



Visible Seam: Window Size



Optimal Window Size



Optimal Window Size

- To avoid seams
 - $\text{window} \geq \text{size of largest prominent feature}$
- To avoid ghosting
 - $\text{window} \leq 2 * \text{size of smallest prominent feature}$
- Depends on the changes (frequencies) in the image
 - hard to find
- We will use image pyramids to avoid the need to find window size

Image Pyramids

What is Blending?

Merge two image in a seamless way





Image Pyramids

- Pyramid of an Image: A collection of the same image, reduced size by half at each level



Size:1x1



Size:1/2x1/2




1/4x1/4



1/8x1/8

Why Pyramids?

Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N = 2^k$)




Gaussian Pyramid [Burt and Adelson, 1983]


Image Pyramids: How?

- Naïve Solution?

Subsampling: take the 2nd pixel from each row/col




Size:1x1



Size:1/2x1/2


Image Pyramids: Sub-sampling

Simple subsampling results in aliasing



Forsyth and Ponce

Aliasing



- When sub-sample, the rate should be adjust to capture all the changes in the signal
- The sampling rate at each level in image pyramids is $\frac{1}{2} \times \frac{1}{2}$
- The sampling rate after 3 levels with respect to the original image is $\frac{1}{8} \times \frac{1}{8}$
- After several levels, aliasing is almost inevitable

Aliasing





Image Pyramids: Sub-sampling

Aliasing means that we can not reconstruct something similar



Aliasing



Gaussian Pyramids



G_0

Fig.2a. The Gaussian pyramid. The original image, G_0 , is repeatedly filtered and subsampled to generate the sequence of reduced resolution image G_1 , G_2 , etc. These comprise a set of lowpass-filtered copies of the original image in which the bandwidth decreases in one-octave steps.



G_1



G_2



G_3



G_4


Gaussian Pyramids: How

- We define the Reduce operation
- Input: Image $N \times M$, Output: Image $N/2 \times M/2$
- The Reduce Operation:
 - Blur
 - Convolve with a 5×5 gaussian filter
 - Sub-sample
 - Select only every 2nd pixel in every 2nd row

Gaussian Smoothing Kernel

The contribution of each pixel for the pixels in the next level

Example: 5 Weights: (a, b, c, b, a)



Conditions:

$c > b > a$ (Unimodal)


$2a + 2b + c = 1$ (Receiving weights)

$c + 2a = 2b$ (Contributing weights)

Commonly Used - Binomial Coefficients

	1	2	1			
	1	4	6	4	1	
1	6	15	20	15	6	1

Reduce Operation



Pixel Contributions




Fig 1. A one-dimensional graphic representation of the process which generates a Gaussian pyramid. Each row of dots represents nodes within a level of the pyramid. The value of each node in the zero level is just the gray level of a corresponding image pixel. The value of each node in a high level is the weighted average of node values in the next lower level. Note that node spacing doubles from level to level, while the same weighting pattern or "generating kernel" is used to generate all levels.

Computational Complexity

- Memory
 - $N \times N (1 + 1/4 + 1/16 + \dots) = 4/3 * N \times N$
- Each level can be computed with a single convolution

Gaussian Pyramid

G0



$G1 = \text{Reduce}(G0)$



$G2 = \text{Reduce}(G1)$





512

256

128

64

32

16

8






Figure from David Forsyth

Expand

- Given a Pyramid level K , can we reconstruct level $K-1$?
- First problem: level $k-1$ is not the same size as level K
- The Expand Operation:
 - Zero Padding ($a_1, 0, a_2, 0, a_3, 0, \dots$)
 - Blur
 - Note: Blur needs different normalizations!
 - What is zero padding followed by blur with $\frac{1}{2}(1,2,1)$

Reconstruct: Expand



Reduce and Expand

G0



G1
Reduce(G0)



Expand(G1)



Pyramid Level Difference



G0



Expand(G1)

Pyramid Level Difference

- $L0 = G0 - \text{Expand}(G1)$



Reconstruction

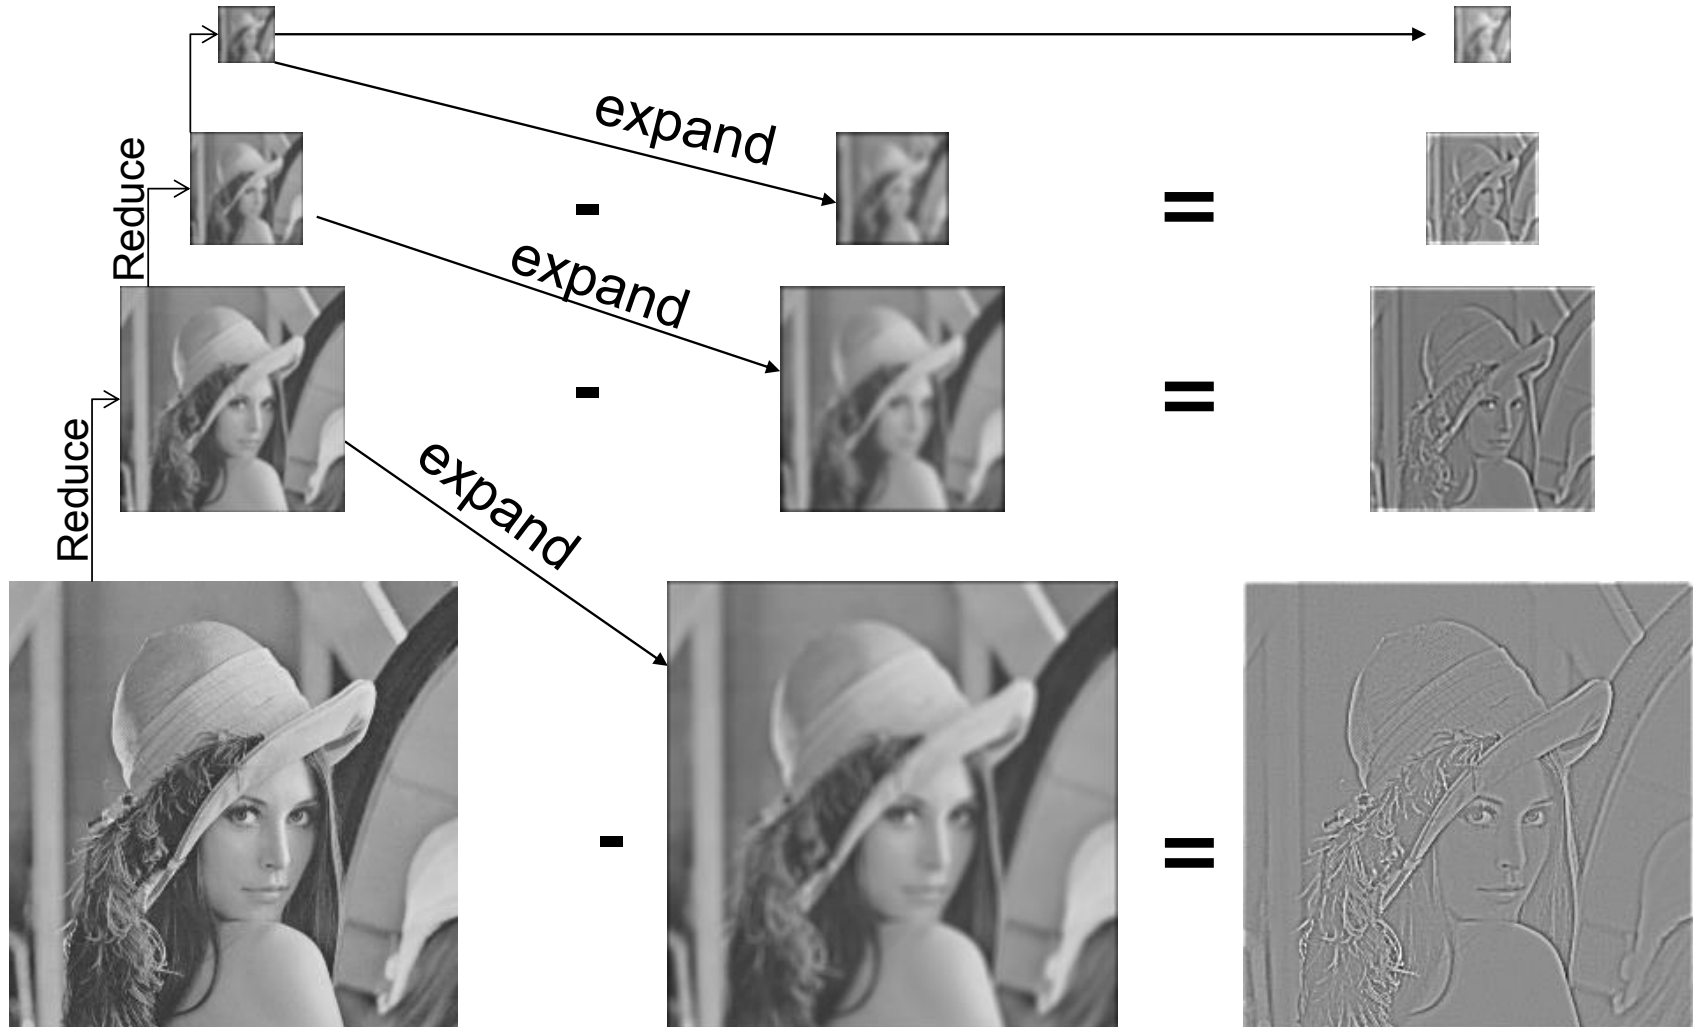
- The difference stores all the details that were lost by the blurring and upsampling operations



Gaussian and Laplacian Pyramid


Gaussian Pyramid

Laplacian Pyramid




Reconstruction: Laplacian and Gaussian pyramids

- The last level in both Laplacian and Gaussian are equal
 - So we can reconstruct the original image
 - How?



Reconstruction

We can reconstruct the original image by storing only the Laplacian pyramid, which is more compressed than the original image



Credit: Shai Avidan

Recall: Optimal Window Size

- To avoid seams
 - window \geq size of largest prominent feature
- To avoid ghosting
 - window $\leq 2 \times$ size of smallest prominent feature

Depends on the changes (frequencies) in the image, hard to find

Pyramid to rescue

- The key idea
 - For a given image, we do not know what are the frequencies in the image
 - But for the pyramid, we can approximate the optimal window size

Pyramid Blending Arbitrary Shape

- Given two images A and B , and a binary image mask M
- Construct Laplacian Pyramids L_a and L_b
- Construct a Gaussian Pyramid G_m
- Create a third Laplacian Pyramid L_c where for each level k

$$L_c(i, j) = G_m(i, j)L_a(i, j) + (1 - G_m(i, j))L_b(i, j)$$

- Reconstruct all levels L_c in to get the blended image

Input



Output

